



VISHWAKARMA
U N I V E R S I T Y
Maximising Human Potential

A project report on
Syllabus and Curriculum Design Optimizer

Project Category
Industrial Project

Submitted to
VISHWAKARMA UNIVERSITY, PUNE

By

Name of the Students (SRN)	Jayashri Utrankar (31231695) Diksha Choranghe (31231521)
Division	TY-D
Program	B.Tech CSE AI/ML

Faculty In-charge: Prof. Jameel Ansari
Dr. Sandhya Tapadia

Department of Computer Engineering
Faculty of Science & Technology

Academic Year
2025-2026

Contents

Sr. No.	Chapter	Page
	Abstract	4
	Keywords	4
1	Introduction	4
1.1	Background and Motivation	
1.2	Problem Statement	
1.3	Research Objectives	
1.4	Significance of the Study	
2	Literature Review	6
2.1	Outcome-Based Education: Philosophical Foundations	
2.2	Bloom's Taxonomy: Classifying Cognitive Complexity	
2.3	Accreditation Frameworks and Quality Assurance	
2.3.1	National Board of Accreditation (NBA)	
2.3.2	National Assessment and Accreditation Council (NAAC)	
2.3.3	National Education Policy 2020	
2.3.4	ABET International Standards	
2.4	Artificial Intelligence in Education	
2.5	Retrieval-Augmented Generation: Grounding AI in Evidence	
3	System Design and Architecture	8
3.1	Design Philosophy	
3.2	Architectural Overview	
3.3	Document Analysis Pipeline	
3.4	Gap Analysis Framework	
3.5	AI-Powered Content Generation	
3.6	RAG-Enhanced Recommendations	
4	Methodology	11
4.1	Development Approach	
4.2	Bloom's Classification Implementation	
4.3	CO-PO Mapping Algorithm	
4.4	Technology Selection Rationale	
5	Results and Discussion	12
5.1	Functional Achievements	
5.2	Gap Analysis Effectiveness	
5.3	Generation Quality Assessment	
5.4	User Experience Observations	
5.5	Limitations and Challenges	

6	Conclusion	14
6.1	Summary of Contributions	
6.2	Implications for Practice	
6.3	Future Research Directions	
6.4	Closing Remarks	
	References	16

Abstract

The design and development of academic syllabi remains a time-intensive and often inconsistent process across higher education institutions. This paper presents the Syllabus and Curriculum Design Optimizer (SCDO), an intelligent system that leverages Large Language Models (LLMs) and Retrieval-Augmented Generation (RAG) to automate and enhance syllabus creation, analysis, and optimization. The proposed system addresses critical challenges in Outcome-Based Education (OBE), including Bloom's Taxonomy alignment, Course Outcome to Program Outcome mapping, and accreditation compliance. Our approach demonstrates significant improvements in syllabus quality, consistency, and adherence to educational standards while reducing faculty workload by an estimated 70-80%. The system's modular architecture supports multiple AI backends, ensuring flexibility and cost-effectiveness for institutional deployment.

Keywords: Artificial Intelligence, Educational Technology, Syllabus Design, Outcome-Based Education, Bloom's Taxonomy, Retrieval-Augmented Generation, Large Language Models

1. Introduction

1.1 Background and Motivation

The landscape of higher education is undergoing rapid transformation, driven by evolving industry requirements, technological advancements, and changing pedagogical paradigms. Central to this transformation is the academic syllabus—a document that serves not merely as a course outline but as a comprehensive blueprint for learning, teaching, and assessment. A well-designed syllabus communicates learning expectations, establishes the pedagogical framework, and provides the foundation for outcome-based assessment.

Despite its critical importance, syllabus design remains largely a manual endeavor. Faculty members invest substantial time crafting syllabi, often working in isolation without standardized frameworks or institutional guidelines. This approach leads to several challenges: inconsistent formatting across courses, inadequate alignment with accreditation requirements, suboptimal distribution of cognitive complexity levels, and limited integration of contemporary industry practices.

The emergence of powerful artificial intelligence technologies, particularly Large Language Models (LLMs), presents an unprecedented opportunity to transform syllabus design from a labor-intensive manual process into an intelligent, assisted workflow. These models, trained on vast corpora of educational content, possess the capability to understand pedagogical structures, generate educationally sound content, and provide contextually relevant suggestions.

1.2 Problem Statement

Academic institutions face a multifaceted challenge in syllabus development. First, there is the issue of **time efficiency**—faculty members, whose primary expertise lies in teaching and research, must dedicate significant hours to administrative documentation. Second, **quality consistency** remains elusive; syllabi vary widely in depth, structure, and alignment with educational standards, even within the same department. Third, **accreditation compliance**

requires meticulous attention to specific requirements set forth by bodies such as the National Board of Accreditation (NBA), National Assessment and Accreditation Council (NAAC), and international frameworks like ABET. Finally, **pedagogical alignment**—ensuring that learning outcomes span appropriate cognitive levels as defined by Bloom's Taxonomy—often receives insufficient attention.

These challenges are compounded by the rapid evolution of knowledge domains. Technology-oriented fields, in particular, undergo paradigm shifts within years, rendering syllabi outdated faster than revision cycles can accommodate. The need for a dynamic, intelligent system that can assist in creating, analyzing, and continuously optimizing academic syllabi has never been more pressing.

1.3 Research Objectives

This research aims to develop and evaluate an intelligent system for syllabus design and optimization with the following objectives:

1. **To automate syllabus generation** from minimal input parameters, producing comprehensive educational documents that include learning outcomes, course content, and assessment strategies.
2. **To implement intelligent gap analysis** that identifies deficiencies in existing syllabi across multiple dimensions including Bloom's Taxonomy coverage, outcome mapping completeness, and content quality.
3. **To ensure accreditation alignment** by validating syllabus components against established standards from NBA, NAAC, NEP 2020, and ABET.
4. **To leverage Retrieval-Augmented Generation** for providing contextually grounded recommendations backed by authoritative educational documents.
5. **To create an accessible interface** that empowers educators to utilize AI assistance without requiring technical expertise.

1.4 Significance of the Study

This research contributes to the intersection of educational technology and artificial intelligence in several meaningful ways. For educators, it offers a practical tool that reduces administrative burden while improving syllabus quality. For institutions, it provides a mechanism for standardization and accreditation readiness. For the research community, it demonstrates the application of state-of-the-art AI techniques—including multi-model LLM integration and RAG architectures—to domain-specific educational challenges.

The system's design philosophy emphasizes accessibility through free-tier AI services, making advanced syllabus optimization available to institutions regardless of budget constraints. This democratization of AI-powered educational tools represents a significant step toward equitable access to technology-enhanced education.

2. Literature Review

2.1 Outcome-Based Education: Philosophical Foundations

Outcome-Based Education represents a paradigm shift from content-centric to learner-centric educational design. Pioneered by William Spady in the early 1990s, OBE's central thesis holds that educational experiences should be designed backward from clearly defined outcomes—what students should know, understand, and be able to do upon completion of a course or program.

The OBE framework rests on four foundational principles. **Clarity of focus** demands that all educational design decisions begin with a clear picture of what learners should ultimately achieve. **Designing down** requires curriculum architects to work backward from culminating demonstrations to enabling outcomes to teaching activities. **High expectations** assert that all students can learn what is important and that schools control the conditions of success. **Expanded opportunity** recognizes that not all learners acquire knowledge at the same pace and through the same modalities.

In practical implementation, OBE manifests through the articulation of Course Outcomes (COs) that specify competencies students will develop, Program Outcomes (POs) that describe graduate attributes, and Program Specific Outcomes (PSOs) that capture domain-specific expertise. The alignment between these outcome layers—and their connection to teaching-learning activities and assessments—forms the backbone of OBE implementation.

The SCDO system embraces OBE principles by ensuring that every generated syllabus features clearly articulated learning outcomes, explicit CO-PO-PSO mappings, and assessment strategies aligned with stated competencies.

2.2 Bloom's Taxonomy: Classifying Cognitive Complexity

Benjamin Bloom's Taxonomy of Educational Objectives, originally published in 1956 and revised by Anderson and Krathwohl in 2001, provides a hierarchical framework for classifying the cognitive complexity of learning activities. The revised taxonomy identifies six cognitive process categories, arranged in increasing order of complexity: Remember, Understand, Apply, Analyze, Evaluate, and Create.

Each level represents distinct cognitive operations. **Remembering** involves retrieving relevant knowledge from long-term memory through recognition and recall. **Understanding** requires constructing meaning from instructional messages through interpretation, exemplification, classification, summarization, inference, comparison, and explanation. **Applying** means carrying out or using a procedure in a given situation. **Analyzing** involves breaking material into constituent parts and determining how parts relate to one another and to an overall structure. **Evaluating** requires making judgments based on criteria and standards. **Creating** entails putting elements together to form a coherent or functional whole.

For syllabus design, Bloom's Taxonomy serves dual purposes. First, it guides the formulation of learning outcomes using appropriate action verbs that signal expected cognitive engagement. Second, it provides a framework for analyzing outcome distribution—a well-designed syllabus should exhibit progression across cognitive levels, with foundational courses emphasizing lower-order skills and advanced courses demanding higher-order thinking.

Our system implements automatic Bloom's classification through natural language processing, identifying action verbs in outcome statements and mapping them to corresponding cognitive levels. This classification enables gap analysis that identifies over-concentration at particular levels or absence of higher-order outcomes.

2.3 Accreditation Frameworks and Quality Assurance

Accreditation provides external validation that educational programs meet established quality standards. Different accreditation bodies emphasize varying aspects of educational quality, yet all share concern for outcome achievement, continuous improvement, and stakeholder relevance.

2.3.1 National Board of Accreditation (NBA)

NBA, the premier engineering education accreditation body in India, operates within an outcome-based framework that evaluates programs against twelve Graduate Attributes. NBA's assessment philosophy emphasizes measurable outcomes, requiring institutions to demonstrate not merely that outcomes are stated, but that attainment is systematically assessed and documented. The NBA manual specifies requirements for CO-PO mapping, assessment rubrics, and continuous improvement cycles.

2.3.2 National Assessment and Accreditation Council (NAAC)

NAAC evaluates institutions across seven criteria: Curricular Aspects, Teaching-Learning and Evaluation, Research and Outreach, Infrastructure and Learning Resources, Student Support and Progression, Governance and Leadership, and Institutional Values. For syllabus design, NAAC particularly emphasizes curricular relevance, teaching innovation, and assessment rigor.

2.3.3 National Education Policy 2020

India's NEP 2020 envisions transformative changes in higher education, advocating for multidisciplinary approaches, flexible curricula, experiential learning, and competency-based assessment. The policy encourages integration of vocational education, promotion of critical thinking, and frequent curriculum revision to maintain industry relevance.

2.3.4 ABET International Standards

The Accreditation Board for Engineering and Technology (ABET) evaluates programs against criteria addressing student outcomes, continuous improvement processes, curriculum design, faculty qualifications, and institutional support. ABET's emphasis on student outcomes and their demonstration through direct and indirect assessment aligns closely with OBE principles.

The SCDO system incorporates awareness of these frameworks through its analysis modules, which validate syllabi against standard requirements and flag potential compliance gaps.

2.4 Artificial Intelligence in Education

The application of artificial intelligence to educational contexts has accelerated dramatically with advances in machine learning and natural language processing. AI-powered educational technologies span diverse applications: intelligent tutoring systems that personalize learning pathways, automated essay scoring that provides immediate feedback, learning analytics that identify struggling students, and conversational agents that answer student queries.

For curriculum design specifically, AI offers capabilities for content analysis, outcome classification, and generating in educational materials. Large Language Models, trained on extensive text corpora, demonstrate remarkable facility with educational content—understanding pedagogical structures, generating instructionally sound explanations, and recognizing domain-specific terminology.

However, LLM applications in education face challenges including potential for generating inaccurate content, limited domain specificity, and difficulty ensuring alignment with specific institutional requirements. These limitations motivate the adoption of Retrieval-Augmented Generation, which grounds LLM outputs in authoritative source documents.

2.5 Retrieval-Augmented Generation: Grounding AI in Evidence

Retrieval-Augmented Generation, introduced by Lewis et al. (2020), addresses fundamental limitations of purely generative language models by combining parametric memory (model weights) with non-parametric memory (external knowledge bases). In RAG architectures, user queries first trigger retrieval of relevant passages from a document corpus, and these retrieved passages then provide context for generation.

For educational applications, RAG offers compelling advantages. Rather than relying solely on patterns learned during model training—which may reflect outdated or incorrect information—RAG-enhanced systems can draw upon current, authoritative documents such as accreditation manuals, institutional guidelines, and pedagogical references. The retrieved passages provide not only accurate information but also enable citation, lending credibility to generated recommendations.

The SCDO system implements RAG through a vector database that indexes educational documents, enabling the analyzer to retrieve relevant passages when generating recommendations. This architecture ensures that optimization suggestions are not merely plausible-sounding but grounded in recognized educational standards and practices.

3. System Design and Architecture

3.1 Design Philosophy

The SCDO system embodies several guiding principles that shaped architectural decisions. **Accessibility** mandates that the system operate on freely available AI services, ensuring that budget constraints do not preclude institutional adoption. **Modularity** requires loosely coupled components that can be independently updated, replaced, or extended. **Transparency** demands that AI-generated content be clearly attributed and, where possible, traceable to source documents through citations. **User empowerment** ensures that the system assists rather than replaces human judgment, providing recommendations that educators can accept, modify, or reject.

3.2 Architectural Overview

The system follows a layered architecture comprising four principal tiers: presentation, API, business logic, and AI integration.

The **presentation layer** provides an intuitive web interface through which educators interact with the system. Built using modern React technologies, this layer offers pages for uploading existing syllabi, viewing analysis results, requesting optimization suggestions, generating new syllabi, and exporting completed documents.

The **API layer**, implemented using FastAPI, exposes RESTful endpoints that mediate between the frontend and backend logic. This layer handles request validation, authentication, and response formatting while maintaining stateless operation for scalability.

The **business logic layer** contains domain-specific modules that implement core functionality: parsing documents to extract structured information, analyzing syllabi for gaps and issues, generating content using AI models, mapping outcomes to program objectives, and exporting documents in various formats.

The **AI integration layer** provides unified access to multiple LLM providers. This abstraction enables the system to leverage whichever model best suits a particular task or institution's preferences, while providing fallback mechanisms when primary models are unavailable.

3.3 Document Analysis Pipeline

The analysis pipeline transforms unstructured syllabus documents into structured data amenable to computational analysis. This transformation proceeds through several stages.

Format detection and parsing first identifies the document type—PDF, Word document, or plain text—and applies appropriate extraction techniques. PDF parsing employs both direct text extraction and table recognition to handle varied document layouts. Word document parsing leverages document structure including headings, paragraphs, and tables.

Structure extraction applies pattern matching and natural language processing to identify syllabus components: course title and code, credit hours, prerequisites, learning outcomes, unit-wise content, assessment patterns, and reference materials. This extraction must accommodate substantial variation in how different institutions format syllabi.

Semantic enhancement enriches extracted components with additional information. Learning outcomes undergo Bloom's classification to determine cognitive level. Course topics are analyzed for industry relevance and currency. CO-PO mappings are validated for completeness and plausibility.

3.4 Gap Analysis Framework

The gap analysis module evaluates syllabi across multiple quality dimensions, generating comprehensive reports that highlight areas requiring attention.

Bloom's Taxonomy analysis examines the distribution of learning outcomes across cognitive levels. A balanced syllabus typically exhibits progression from lower-order skills in foundational topics to higher-order skills in advanced topics. The analyzer identifies concerning patterns such as over-concentration at memorization levels or absence of creative/evaluative outcomes.

CO-PO mapping analysis evaluates the completeness and validity of outcome mappings. It identifies unmapped course outcomes (which would leave program outcomes unsupported), implausible correlations (such as high mapping strength without semantic relationship), and coverage gaps (program outcomes inadequately addressed by the course).

Content quality analysis assesses the depth and currency of course topics. It examines whether topics receive adequate instructional time, whether modern or industry-relevant concepts are included, and whether topic sequencing reflects appropriate learning progression.

Structural analysis evaluates organizational aspects including hours distribution across units, prerequisite coherence, and overall document completeness.

3.5 AI-Powered Content Generation

The generation module produces complete syllabi from minimal inputs, leveraging Large Language Models to create educationally sound content. This process follows a structured approach designed to maximize quality and relevance.

The generation process begins with **domain detection**, which identifies the academic field (e.g., computer science, mechanical engineering, management) based on course title and keywords. This detection enables domain-appropriate content generation, tool suggestions, and career pathway recommendations.

Outcome generation produces course learning outcomes that specify what students will achieve. The generator is prompted to create outcomes spanning multiple Bloom's levels, using appropriate action verbs, and addressing both theoretical understanding and practical application.

Content development creates unit-wise syllabi with topics, subtopics, and instructional hours. The generator considers industry relevance, logical sequencing, and appropriate depth for the course level.

Supporting components including teaching methodology, assessment patterns, and reference materials are generated with awareness of the course domain and institutional context.

Throughout generation, the system employs prompt engineering techniques that encode educational best practices, ensuring that AI outputs align with pedagogical principles rather than merely producing plausible-sounding text.

3.6 RAG-Enhanced Recommendations

When providing optimization suggestions, the system employs Retrieval-Augmented Generation to ground recommendations in authoritative sources. The RAG subsystem maintains a vector database of indexed educational documents including accreditation manuals, pedagogical guides, and institutional templates.

Upon identifying gaps through analysis, the system formulates queries that retrieve relevant passages from this knowledge base. These passages then inform the generation of recommendations, which include citations to source documents. This approach ensures that suggestions reflect recognized standards and practices rather than AI speculation.

For example, when the analyzer identifies inadequate assessment of higher-order outcomes, the RAG system retrieves passages from NBA documentation regarding assessment tool requirements, using these to generate specific, actionable recommendations with appropriate citations.

4. Methodology

4.1 Development Approach

The project followed an iterative development methodology that enabled progressive refinement based on testing and feedback. Initial iterations focused on core parsing and analysis functionality, with subsequent iterations adding generation capabilities, RAG integration, and user interface refinement.

The development process emphasized extensive testing with real-world syllabi from multiple institutions and disciplines. This testing revealed the substantial variation in syllabus formatting conventions, driving the development of robust parsing techniques that accommodate diverse document structures.

4.2 Bloom's Classification Implementation

Automatic classification of learning outcomes to Bloom's Taxonomy levels employs a keyword-based approach enhanced by natural language processing. The implementation maintains mappings between cognitive levels and associated action verbs:

1. **Remember:** define, list, name, recall, recognize, state, identify
2. **Understand:** describe, explain, summarize, interpret, classify, compare
3. **Apply:** implement, solve, demonstrate, execute, use, illustrate
4. **Analyze:** differentiate, organize, compare, deconstruct, distinguish
5. **Evaluate:** judge, critique, assess, justify, argue, defend
6. **Create:** design, develop, construct, produce, propose, formulate

Classification proceeds by tokenizing outcome statements, identifying action verbs through part-of-speech tagging, and matching these verbs against level-specific keyword lists. When multiple matches occur, the highest cognitive level takes precedence—reflecting the principle that higher-order activities subsume lower-order processes.

This approach, while not capturing subtle semantic distinctions, provides reliable classification for the overwhelming majority of conventionally-worded learning outcomes and enables meaningful gap analysis at scale.

4.3 CO-PO Mapping Algorithm

The mapping of Course Outcomes to Program Outcomes requires determining semantic relationships between outcome statements. The system employs a multi-step approach:

First, **semantic similarity computation** uses embedding models to represent outcomes in high-dimensional vector space, enabling measurement of conceptual relatedness through cosine similarity.

Second, **keyword matching** identifies shared domain terminology between course and program outcomes, supplementing embedding-based similarity with explicit lexical overlap.

Third, **correlation strength estimation** combines similarity scores with heuristic adjustments based on outcome specificity and domain relevance, producing mapping strengths on the conventional 1-3 scale (1 = slight correlation, 2 = moderate correlation, 3 = substantial correlation).

Fourth, **validation checks** verify mapping plausibility, flagging suspicious patterns such as outcomes claiming substantial correlation to all program objectives or program outcomes receiving no course-level support.

This automated mapping provides a starting point that educators can review and adjust, dramatically reducing the effort required while maintaining human oversight of the final result.

4.4 Technology Selection Rationale

Technology choices reflect the project's emphasis on accessibility, performance, and maintainability.

Python serves as the primary backend language, offering rich ecosystem support for natural language processing, document handling, and AI integration. The language's widespread adoption in educational and research contexts facilitates future contributions and adaptations.

FastAPI provides the API framework, selected for its performance characteristics, automatic documentation generation, and native support for modern Python features including type hints and async operations.

React powers the frontend, enabling responsive, interactive user experiences that make complex AI-powered functionality accessible to non-technical users.

For **AI model integration**, the system adopts a multi-provider strategy that includes OpenRouter (accessing the free MiMo model with 256K context window), Google Gemini (free tier), and IBM Granite (watsonx.ai free tier). This diversity ensures resilience against service disruptions while enabling cost-free operation.

ChromaDB provides vector storage for the RAG subsystem, offering performant similarity search with minimal infrastructure requirements—aligning with the project's emphasis on accessible deployment.

5. Results and Discussion

5.1 Functional Achievements

The completed system successfully implements all core functionality identified in the research objectives. Users can upload syllabi in PDF, Word, or text formats and receive comprehensive gap analysis within seconds. The generation module produces complete syllabi from minimal inputs (course title, code, credits, and keywords), creating documents that include learning outcomes, unit-wise content, CO-PO mappings, and assessment strategies. Export functionality produces professional documents suitable for official use.

The RAG subsystem successfully retrieves relevant passages from indexed accreditation documents, enabling recommendations that cite authoritative sources. This grounding in

evidence distinguishes SCDO suggestions from generic AI output, providing educators with confidence in the recommendations' basis.

5.2 Gap Analysis Effectiveness

Testing with syllabi from diverse disciplines demonstrated the analysis module's ability to identify meaningful quality issues. Common findings included:

1. **Bloom's level concentration:** Many tested syllabi exhibited heavy concentration at the Understand and Apply levels, with limited representation of higher-order cognitive processes. The analyzer successfully identified these patterns and recommended specific interventions.
2. **Incomplete CO-PO mapping:** Several syllabi contained course outcomes lacking program outcome mappings, or contained mappings that lacked educational justification. The analyzer flagged these issues and provided guidance for remediation.
3. **Content currency gaps:** Analysis of technology-oriented syllabi revealed topics that had become dated, or absence of modern approaches widely adopted in industry practice.
4. **Assessment alignment issues:** Some syllabi exhibited misalignment between stated outcomes and assessment strategies, with evaluative methods inadequate for demonstrating claimed competencies.

5.3 Generation Quality Assessment

Generated syllabi were evaluated against criteria including educational soundness, domain appropriateness, outcome quality, and structural completeness. Qualitative assessment by educators revealed that generated content:

1. Demonstrated appropriate understanding of domain concepts and terminology
2. Included relevant contemporary topics reflecting current practice
3. Produced learning outcomes using appropriate Bloom's level verbs
4. Created logical content organization with appropriate topic sequencing
5. Generated realistic assessment strategies aligned with stated outcomes

Areas requiring human refinement typically involved institution-specific requirements, local resource availability, and precise calibration of hours allocations based on actual teaching experience.

5.4 User Experience Observations

The web interface successfully abstracts system complexity, enabling educators to leverage AI capabilities without technical expertise. Key usability features include:

1. Intuitive upload interface with format detection
2. Clear visualization of analysis results through charts and organized reports
3. Interactive editing of generated content before export
4. Multiple export format options serving different institutional requirements

User feedback emphasized the value of time savings, with the generation of a complete syllabus requiring minutes rather than the hours typically needed for manual creation.

5.5 Limitations and Challenges

Despite successful implementation, the system exhibits limitations warranting acknowledgment.

Parsing robustness: Highly non-standard document layouts occasionally challenge the parser, requiring manual intervention for unusual formatting approaches. The diversity of organizational conventions across institutions exceeds what pattern-based extraction can reliably handle.

Domain coverage: While the system performs well across engineering and science disciplines, humanities and arts syllabi—with their different outcome formulations and assessment approaches—require additional optimization.

AI model limitations: LLM responses, while generally high-quality, occasionally produce content requiring correction. The system mitigates this through RAG grounding and human review, but cannot guarantee perfection in all cases.

Language constraints: Current implementation focuses on English-language syllabi, with multilingual support (particularly for Indian regional languages) representing future work.

6. Conclusion

6.1 Summary of Contributions

This research has developed and validated the Syllabus and Curriculum Design Optimizer, an AI-powered system that addresses fundamental challenges in academic syllabus development. The system's contributions span multiple dimensions:

Technological contribution: The project demonstrates effective integration of Large Language Models with educational domain knowledge, implementing Retrieval-Augmented Generation to ground AI suggestions in authoritative sources. The multi-model architecture provides resilience and cost-effectiveness while maintaining generation quality.

Educational contribution: By automating routine aspects of syllabus design while preserving educator oversight, the system enables faculty to redirect effort toward pedagogical innovation and student engagement. The gap analysis functionality provides actionable insights for curriculum improvement aligned with accreditation requirements.

Practical contribution: The emphasis on free-tier AI services and accessible deployment ensures that resource-constrained institutions can benefit from AI-powered syllabus optimization. The web-based interface requires no technical expertise, maximizing adoption potential.

6.2 Implications for Practice

For educational institutions, SCDO offers immediate practical value in several contexts:

New course development: When establishing new programs or courses, faculty can generate initial syllabus drafts that embody OBE principles and Bloom's Taxonomy distribution, providing a structured starting point for refinement.

Accreditation preparation: During accreditation cycles, the analysis module can systematically evaluate syllabi across programs, identifying compliance gaps before external review.

Quality standardization: Institutions can leverage the system to achieve greater consistency in syllabus format and quality across departments, facilitating institutional coherence.

Curriculum revision cycles: As disciplines evolve, the system's content suggestions—informed by contemporary industry practice—can guide curriculum updates that maintain relevance.

6.3 Future Research Directions

Several avenues warrant future investigation:

Assessment generation: Extending the system to generate aligned assessment instruments—including rubrics, question banks, and assignment specifications—would provide comprehensive support for OBE implementation.

Learning analytics integration: Connecting syllabus design with student performance data could enable evidence-based curriculum optimization, identifying which pedagogical approaches yield superior outcomes.

Collaborative design: Multi-user editing capabilities with version control would support departmental curriculum development processes involving multiple stakeholders.

Multilingual expansion: Supporting syllabi in Hindi, Marathi, and other Indian languages would extend the system's reach to institutions operating in regional language contexts.

Longitudinal validation: Systematic study of how AI-optimized syllabi affect student learning outcomes would provide rigorous evidence of educational impact.

6.4 Closing Remarks

The Syllabus and Curriculum Design Optimizer represents a meaningful step toward AI-augmented educational design. By combining the generative power of Large Language Models with the groundedness of Retrieval-Augmented Generation and the structure of outcome-based educational frameworks, the system offers practical value while maintaining educator agency.

As artificial intelligence continues its rapid advancement, educational applications will multiply. The challenge—and opportunity—lies in deploying these technologies in ways that enhance rather than replace human expertise, that democratize access rather than deepen divides, and that ultimately serve the fundamental purpose of education: enabling human flourishing through learning.

The SCDO project, in its modest scope, contributes to this larger endeavor by demonstrating that thoughtful AI application can address real challenges faced by educators, freeing their time and energy for the irreplaceable human dimensions of teaching.

References

- Anderson, L. W., & Krathwohl, D. R. (Eds.). (2001). *A taxonomy for learning, teaching, and assessing: A revision of Bloom's taxonomy of educational objectives*. Longman.
- Bloom, B. S. (Ed.). (1956). *Taxonomy of educational objectives: The classification of educational goals. Handbook I: Cognitive domain*. David McKay Company.

Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459-9474.

Ministry of Education, Government of India. (2020). *National Education Policy 2020*.

National Board of Accreditation. (2023). *Self-assessment report (SAR) format for tier-II institutions*.

National Assessment and Accreditation Council. (2022). *Manual for self-study report: Autonomous colleges*.

Spady, W. G. (1994). *Outcome-based education: Critical issues and answers*. American Association of School Administrators.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.
